Conflux V2 accuracy at 900 simulated seconds

Uniform hardware · IID
Moshpit: 56.37 ± 0.76%

Reducer cap \ Number of chunks	8	16	32	64	128	256
64	—	—	—	58.48 ±0.70	59.28 ±0.74	59.27 ±0.75
32	—	—	55.56 ±0.98	57.28 ±0.87	57.74 ±0.67	57.09 ±0.85
16	—	54.69 ±0.72	55.09 ±0.89	54.97 ±0.89	54.93 ±0.82	54.59 ±0.68
8	49.57 ±0.68	48.97 ±0.59	48.79 ±0.69	48.49 ±0.45	48.21 ±0.52	47.61 ±0.66

Heterogeneous hardware · IID
Moshpit: 44.83 ± 0.53%

Reducer cap \ Number of chunks	8	16	32	64	128	256
64	—	—	—	50.51 ±0.54	49.34 ±0.67	49.39 ±0.97
32	—	—	49.37 ±0.51	50.51 ±0.54	49.34 ±0.67	50.37 ±0.57
16	—	47.70 ±0.74	49.37 ±0.51	50.51 ±0.54	50.85 ±0.49	50.63 ±0.44
8	46.66 ±0.58	48.49 ±0.51	49.63 ±0.70	50.71 ±0.39	50.82 ±0.53	50.50 ±0.53

Uniform hardware · Non-IID (α=0.1)
Moshpit: 40.50 ± 0.55%

Reducer cap \ Number of chunks	8	16	32	64	128	256
64	—	—	—	44.75 ±2.13	45.77 ±1.26	45.80 ±1.24
32	—	—	42.92 ±1.85	43.90 ±2.34	44.07 ±1.22	43.53 ±1.93
16	—	40.55 ±1.86	41.84 ±1.78	42.04 ±1.48	41.37 ±1.87	41.19 ±1.12
8	35.51 ±2.87	36.64 ±2.02	36.41 ±1.72	35.81 ±2.08	35.83 ±1.38	35.36 ±1.94

Heterogeneous hardware · Non-IID (α=0.1)
Moshpit: 29.04 ± 2.48%

Reducer cap \ Number of chunks	8	16	32	64	128	256
64	—	—	—	37.28 ±1.96	36.05 ±2.51	36.19 ±1.93
32	—	—	35.87 ±2.73	37.28 ±1.96	36.05 ±2.51	36.91 ±2.47
16	—	34.85 ±1.56	35.87 ±2.73	37.28 ±1.96	37.68 ±1.81	37.44 ±1.71
8	34.44 ±2.10	35.59 ±1.89	35.86 ±3.03	37.54 ±1.72	37.78 ±1.63	37.29 ±1.97

Final all-node test accuracy (%)

64 nodes · 16 V2 trainers · CIFAR-10 / GNLeNet · three paired seeds (mean ± sample SD)
— / hatched: omitted because reducer cap > chunks; these cells are redundant, not failed runs.
A reducer cap can exceed the number actually used on heterogeneous hardware.